



Parallel Sonar Beam Tracing

CS566 Final Project

Alessandro Febretti, Homam Abu Saleem

ENDURANCE

- **Environmentally Non - Disturbing Under - ice Robotic ANtarctic Explorer**
- Funded by NASA ASTEP program





ENDURANCE Lake Bonney
Drop Sonde Missions
Status: November 18, 2009

- 2008 sonde casts
- 2009 completed
- 2009 new targets

True North

Observer #1
Nov 17, 2009


Observer #2
Nov 18, 2009

GO fly zone

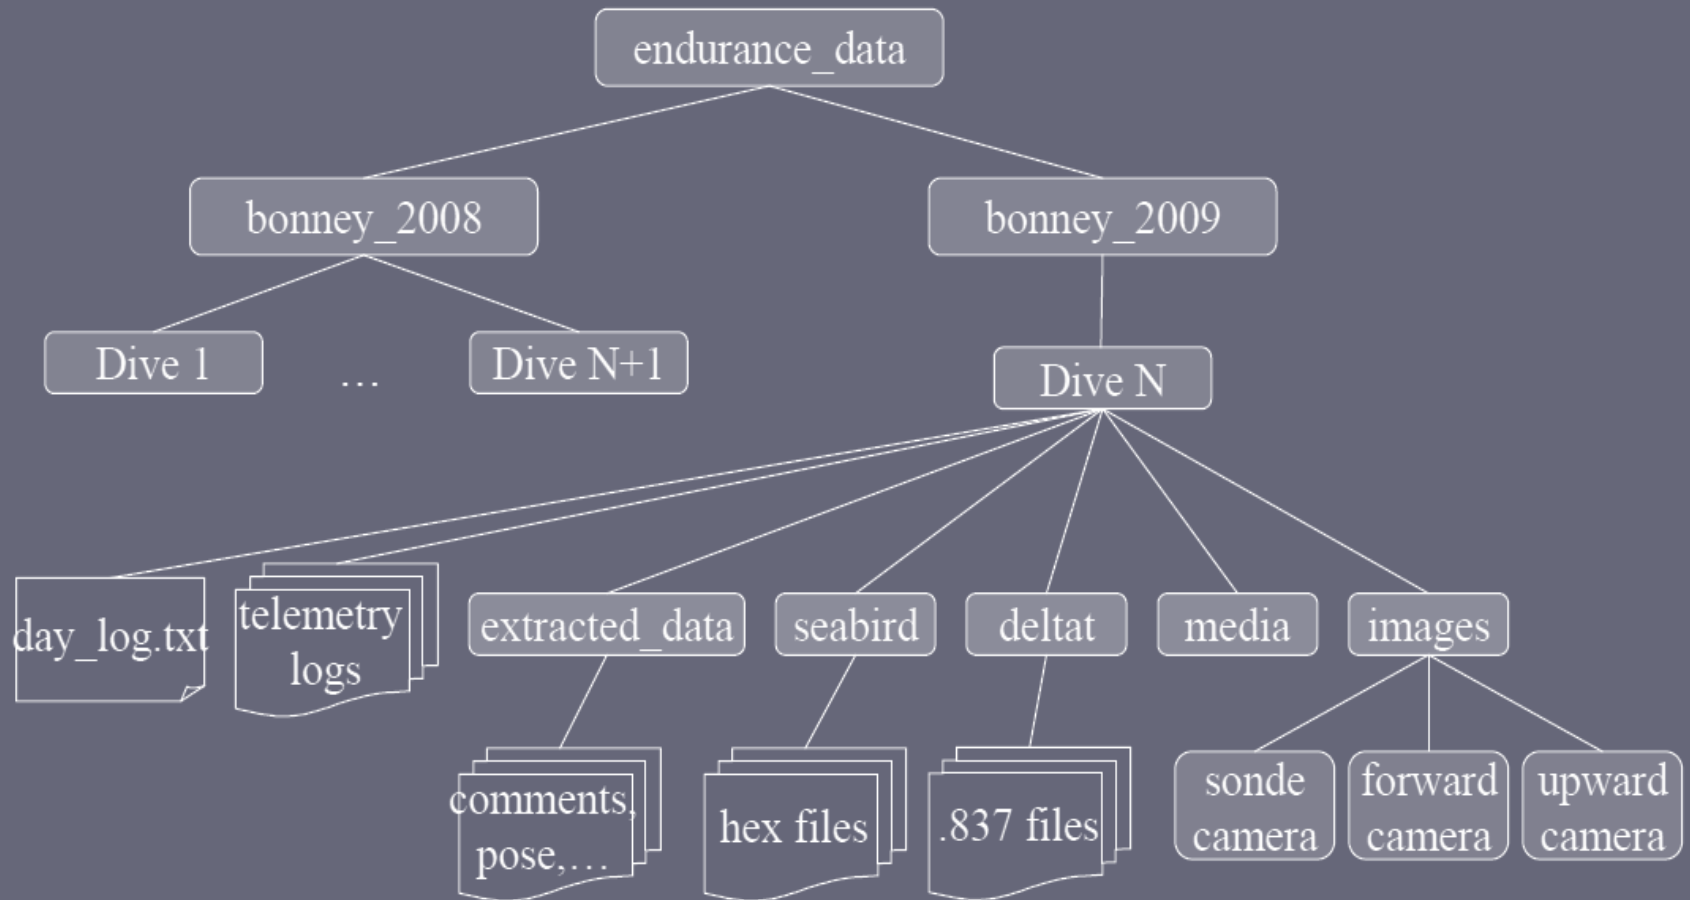
GO fly zone

100 meter grids

Data collection: Sonar



Dataset structure



Objectives

- Create a new, high resolution lake bathymetry
 - Compute precise lake volume
 - Simulate lake interaction with the McMurdo Valley System
- Generate a 3D model of the lake/glacier interface

Challenges

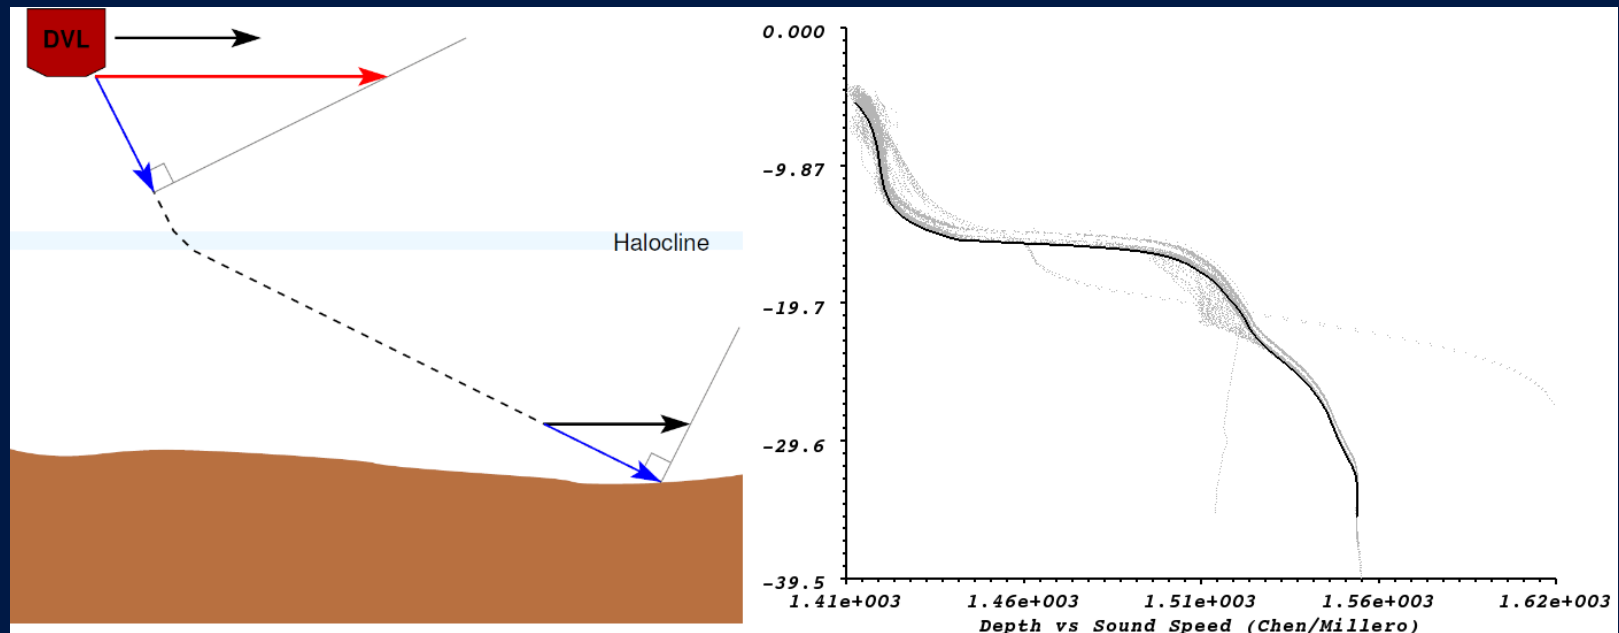
- Sonar data is very imprecise
 - Multipath reflections
 - Water turbidity
 - Navigation / attitude errors
 - Water level variations
 - Lake Chemistry

How Sound Travels Through Water

- Basic measurement: return time of sound pressure wave
 - Converted to distance using estimated **speed of sound**
- Speed of sound in sea water: $\sim 1500\text{m/s}$
 - estimate changes with temperature, salinity, pressure.
 - **Sound speed influences beam paths**
 - Refraction (Snell's Law)

How Sound Travels Through Water

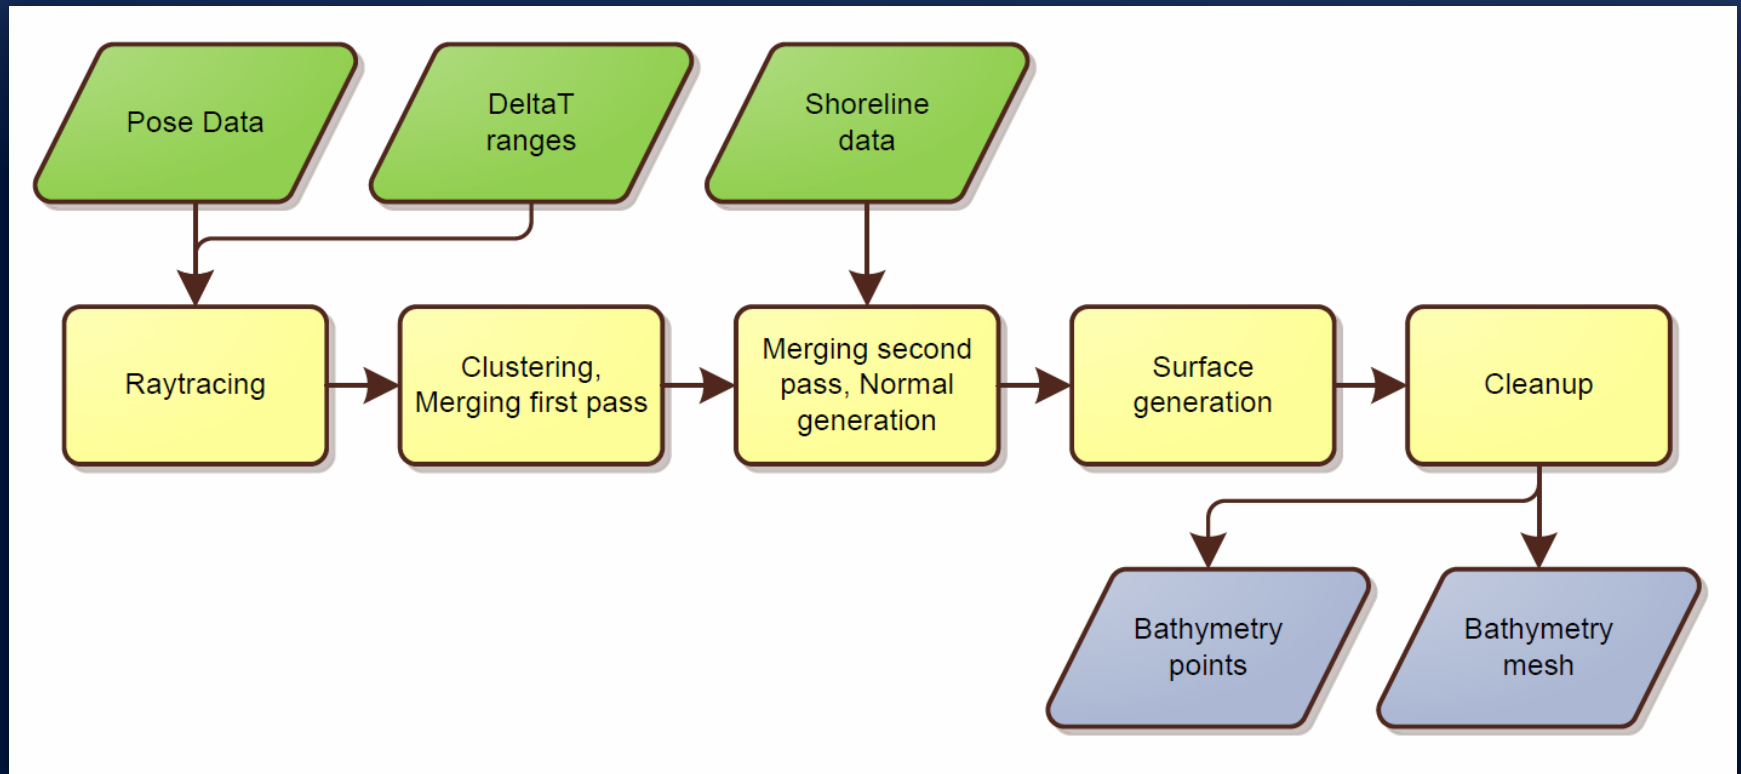
- Different sound velocity models (i.e. DelGrosso, Chen-Millero)
 - Tested in ocean water only



Requirements

- To address objective researchers need to
 - Tweak SVM and beam tracing parameters
 - Apply different navigation corrections
 - Change noise filtering levels
- **Big parameter space**
 - Need fast tweak-process-visualize pipeline
 - Point cloud is enough
 - Surface reconstruction can be done offline (ex. using Poisson reconstruction)

Pipeline



dttools

- Toolkit used for ENDURANCE sonar processing
 - dtrt
 - beam tracer, adapted from MBSystem
 - dtmerge
 - merge, cleanup, normal estimation
 - dtpoisson
 - surface reconstruction
 - dtsample
 - dataset query tool
- We add **mpidtrt**, **mpidtmerge**

Ray-tracing Steps

- Load the configuration files
- Load DeltaT files
- Load Position files
- Merge DeltaT points and Position points
- Trim off some noise points in the DeltaT points
- Calculate distance using sound velocity property, orientation, and position configurations for DeltaTpoints / P
- Generate 3D coordinates using the distance, position and orientation configurations
- Write the coordinates to the output file

Parallel Ray-tracing Steps

- Load the configuration files
- Load the entire DeltaT file to get the data set size
- Load the entire Position file because all of the points are needed for each point in DeltaT
- Each process merges (total DeltaT points / P) and Position points
- Last process does the remainder of DeltaTpoints / P since division will leave a remainder $< P$
- Trim off some noise points in the DeltaT points
- Calculate distance using sound velocity property, orientation, and position configurations for DeltaTpoints / P
- Generate 3D coordinates using the distance, position and orientation configurations DeltaTpoints / P
- Each process writes the coordinates to a separate output file
- Process 0 merges all of the separate output files into one and deletes the separate output file

Merging DeltaT and Position

- Each DeltaT ping has a time stamp
- Each Position data has a time stamp
- Search the time stamp for the positions and pick the closes one to the time stamp of DeltaT ping
- The position data will be used as position data for the DeltaT ping

Removing Noise

- Each ping has multiple sonar signals sent and receive the reflection
- If a received signal is not within a set filter angle, disregard it
- If a received signal has time that is less than a certain threshold, we disregard it

Parallel Ray-tracing Notes

- At least one process has to load the entire DeltaT and Position files to get the size of data since it is unknown, disadvantage
- The writing part has to be sequential, because when multiple processes write to the same file, there will be concurrency problem and some process overwrites another process output
- The sound velocity was calculated by scientists and provided along with the data. It was calculated according to the water profile for each region

Ray-trace Input/Output Samples

- DeltaT: 1259565829.523583 0.000 0.000 0.000 0.000 0.000 0.000 0.000 0.000 0.000 0.000 0.000 0.000 0.000 0.000

0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	0.000	0.000	0.000	39.360	39.520	0.000	39.360	39.200	39.040	39.040	38.880	38.720	38.560	38.400	
38.080	38.080	38.240	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	34.080	34.080	0.000	34.080	33.920	33.920	33.920	33.760	33.760	33.760	33.600	33.600	33.600	33.600	33.440
33.440	33.440	33.280	33.120	33.120	33.120	33.120	32.960	32.960	32.800	32.800	32.800	32.800	32.800	32.640	0.000
32.640	32.480	32.480	32.320	32.320	32.160	32.160	32.160	32.000	32.000	32.000	32.000	32.160	0.000	32.160	
32.000	32.000	32.000	31.840	31.840	31.840	31.680	31.680	31.680	31.520	31.520	31.520	31.520	31.360	31.360	
31.360	31.200	31.200	31.200	31.360	31.360	31.520	31.520	31.680	31.680	31.680	31.680	31.680	31.680	31.520	
31.520	31.520	31.520	31.520	31.520	31.360	31.360	31.040	31.040	32.000	0.000	0.000	0.000	0.000	0.000	
0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	35.520	35.840	35.840
36.000	36.000	36.640	36.640	36.640	36.640	36.480	36.480	36.480	36.480	36.480	0.000	0.000	0.000	0.000	
0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	0.000	0.000	0.000	45.440	45.600	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
- Position: 1259565829.483745 0.001 0.003 5.548 -3.172 -1.014 5.378
- Output: 0.054576, 0.113037, 62.693869


mpidtmmerge: steps

- compute bounds
- merge point set into local grid
- distribute bins
- compute final points
- write points



compute bounds

- Each node works on subset of points
- MPI_AllReduce to compute global min/max

P1



P2




AllReduce


merge point set into local grid

- Each node works on same subset of points again
- global min/max used to create binning grid
 - use octree lookup to quickly find target bin for points
- Accumulate position, normal, num. points for each bin

P1




P2




distribute bins

- Each node is assigned a subset of bins
- bins are exchanged and merged
- each node now has global accumulated positions, normal, point counts for bins it owns

P1




P2




compute final points

- Each node computes final points, discards bins under threshold.
- each node now has part of final solution
- Normal estimation optionally happens here

P1




P2





write output

- Each node appends its part of solution to output file
- Sequential. First node overwrites. Other nodes append.

P1

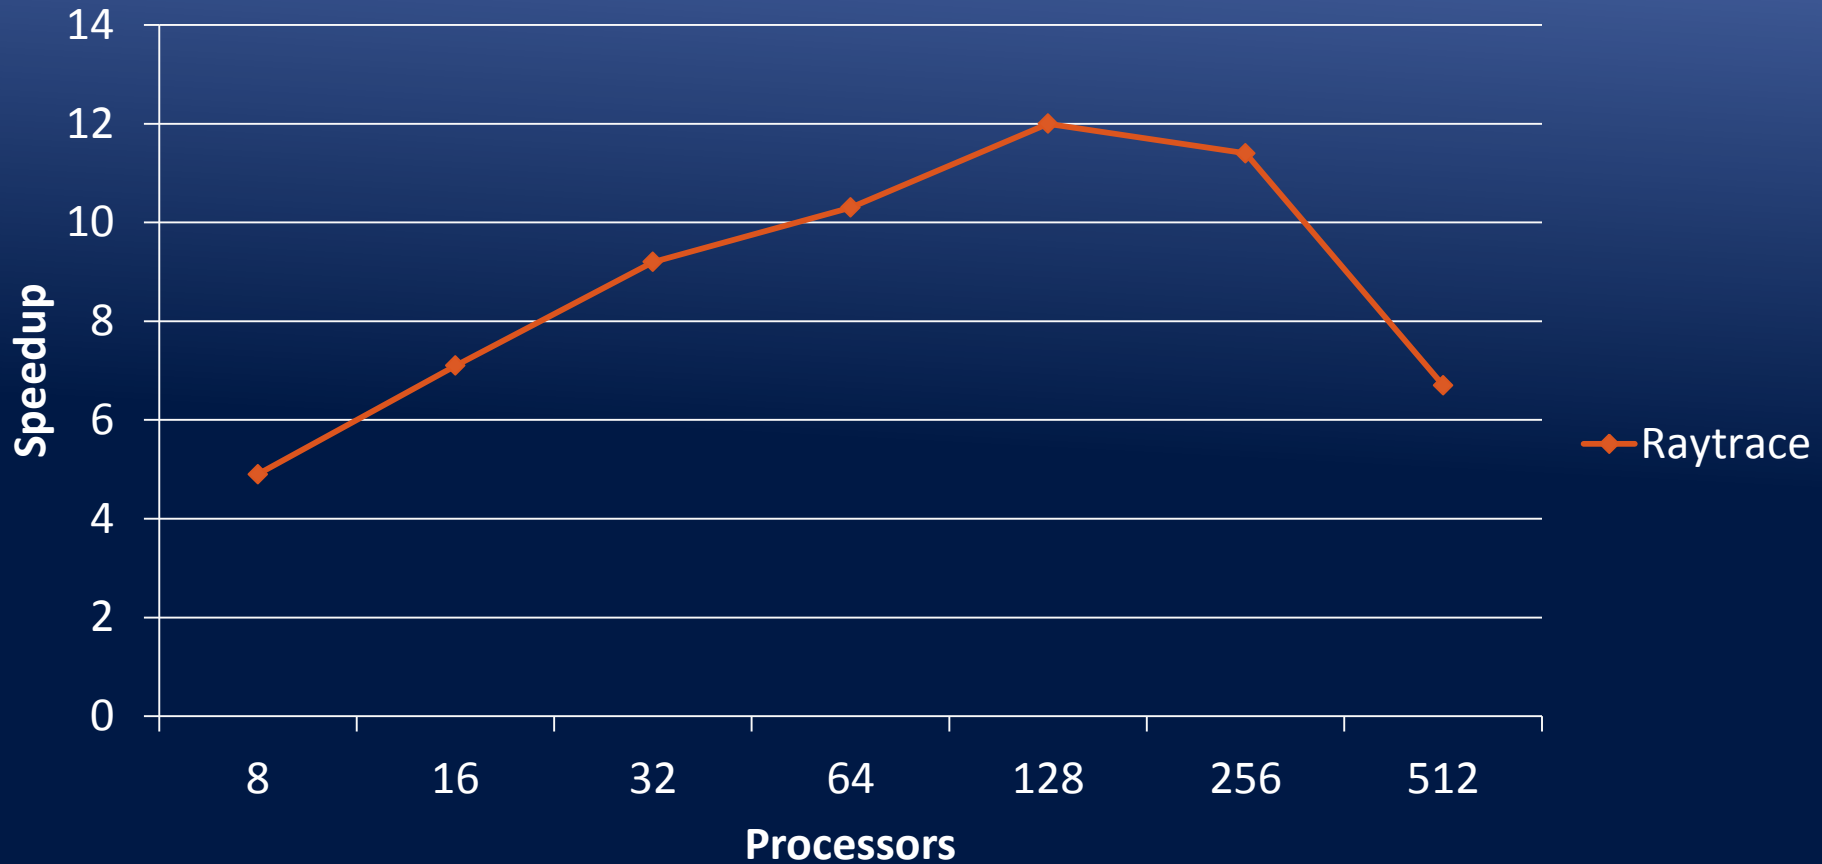


P2




output.xyz

Raytrace



- Run on EVL Lyra cluster (36 machines, 16 cores)
- Best speedup (P=128): **128x**
- Sequential time: **~45.5 minutes**. Best parallel time: **~3.9 minutes**
- **This speedup is for processing each ping of the data set, most of the time we do not need to process each ping, we process something like 1 ping out of 10 pings, so this time would be much less than the above**

Merge Speedup for 1m and 2m voxel size




- Run on EVL Lyra cluster (36 machines, 16 cores)
- Best speedup (P=128): 25-40x
- Sequential time: ~3.5 minutes. Best parallel time: ~5 seconds
- Entire dataset (~200Mil points) can be reprocessed almost in real time

Visualization



Visualization

- xyz data copied to GPU memory as-is
- Geometry shader generates user facing quads from points
- Fragment shader rasterizes sphere within quad
- Scales to a few mil points
- Future work: LOD / spatial optimizations



Future work

- Import full, unmerged data into visualization
- Fast query of visualized data
- Mark regions as noisy, feed back to merge step, reprocess on-the-fly
- Integrate with other data sources
 - Chemistry readings
 - AUV position
 - Map